

## Model Optimization and Tuning Phase

Date	07 July 2024
Team ID	739665
Project Title	BlueBerry Yield Prediction
Maximum Marks	6 Marks

### Hyperparameter Tuning Documentation :

Hyperparameter tuning involves adjusting the parameters that govern the training process of machine learning models to optimize their performance. It includes methods such as grid search, random search, and Bayesian optimization. Proper documentation helps in understanding the impact of different hyperparameters, streamlining the tuning process, and replicating results. Clear records of hyperparameter settings and their outcomes are essential for achieving the best model accuracy and efficiency.

Model	Tuned Hyperparameters	Optimal Values

## Linear Regression

```
from sklearn.linear_model import Ridge
ridge = Ridge()
parameters = {'alpha': [0.1, 1, 10]} # Example values for regularization strength

ridge_regressor = GridSearchCV(ridge, parameters, scoring='neg_mean_squared_error', cv=5)
ridge_regressor.fit(x_train, y_train)

best_alpha = ridge_regressor.best_params_['alpha']
print("Best Alpha:", best_alpha)

# Using the best model found by GridSearchCV
best_ridge = ridge_regressor.best_estimator_
best_ridge.fit(x_train, y_train)
pred_ridge = best_ridge.predict(x_test)
```

```
mae_ridge = mean_absolute_error(y_test, pred_ridge)
mse_ridge = mean_squared_error(y_test, pred_ridge)
rmse_ridge = np.sqrt(mse_ridge)
rsq_ridge = r2_score(y_test, pred_ridge)

print("MAE: %.3f" % mae_ridge)
print("MSE: %.3f" % mse_ridge)
print("RMSE: %.3f" % rmse_ridge)
print("R-Square: %.3f" % rsq_ridge)
print("Training Accuracy:", best_ridge.score(x_train, y_train))
print("Testing Accuracy:", best_ridge.score(x_test, y_test))
```

Best Alpha: 0.1  
MAE: 95.466  
MSE: 14043.502  
RMSE: 118.505  
R-Square: 0.991  
Training Accuracy: 0.991011446378135  
Testing Accuracy: 0.9913088598782471



## RandomForest Regressor

```
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

rf = RandomForestRegressor(random_state=42)

grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, n_jobs=-1, verbose=2)

grid_search.fit(x_train, y_train)

best_params = grid_search.best_params_
best_score = grid_search.best_score_

print(f"Best Parameters: {best_params}")
print(f"Best Cross-Validation Score: {best_score:.3f}")

# Train the model with the best parameters
best_rf = grid_search.best_estimator_
pred_rf_train_tu = best_rf.predict(x_train)
pred_rf_tu = best_rf.predict(x_test)
```

```
mae_rf_train_tu = mean_absolute_error(y_train, pred_rf_train_tu)
mae_rf_tu = mean_absolute_error(y_test, pred_rf_tu)
mse_rf_tu = mean_squared_error(y_test, pred_rf_tu)
rmse_rf_tu = np.sqrt(mse_rf_tu)
rsq_rf_tu = r2_score(y_test, pred_rf_tu)
```

```
print("MAE_train: %.3f" % mae_rf_train_tu)
print("MAE: %.3f" % mae_rf_tu)
print("MSE: %.3f" % mse_rf_tu)
print("RMSE: %.3f" % rmse_rf_tu)
print("R-Square: %.3f" % rsq_rf_tu)
print("Training Accuracy: %.3f" % best_rf.score(x_train, y_train))
print("Testing Accuracy: %.3f" % best_rf.score(x_test, y_test))
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits  
Best Parameters: {'bootstrap': True, 'max\_depth': None, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200}  
Best Cross-Validation Score: 0.986  
MAE\_train: 41.448  
MAE: 110.332  
MSE: 19188.170  
RMSE: 138.521  
R-Square: 0.988  
Training Accuracy: 0.998  
Testing Accuracy: 0.988

# DecisionTree

## Regressor

```
dt = DecisionTreeRegressor()

param_grid = {
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10, 15],
    'min_samples_leaf': [1, 2, 5, 10],
    'max_features': ['auto', 'sqrt', 'log2', None]
}

grid_search = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1)

grid_search.fit(x_train, y_train)

print("Best Parameters:", grid_search.best_params_)
print("Best CV Score:", grid_search.best_score_)

best_dt = grid_search.best_estimator_
pred_dt_tu = best_dt.predict(x_test)
```

```
mae_dt_tu = mean_absolute_error(y_test, pred_dt_tu)
mse_dt_tu = mean_squared_error(y_test, pred_dt_tu)
rmse_dt_tu = np.sqrt(mse_dt_tu)
rsq_dt_tu = r2_score(y_test, pred_dt_tu)
```

```
print("MAE:", mae_dt_tu)
print("MSE:", mse_dt_tu)
print("RMSE:", rmse_dt_tu)
print("R-Squared:", rsq_dt_tu)
print("Training Accuracy:", best_dt.score(x_train, y_train))
print("Testing Accuracy:", best_dt.score(x_test, y_test))
```

Best Parameters: {'max\_depth': None, 'max\_features': None, 'min\_samples\_leaf': 5, 'min\_samples\_split': 10}  
Best CV Score: -40740.29928310072  
MAE: 128.17739583664462  
MSE: 30284.679955869266  
RMSE: 174.02494061446845  
R-Squared: 0.9812576374711801  
Training Accuracy: 0.9931849259250838  
Testing Accuracy: 0.9812576374711801

# XGBoost

## Regressor

```
xgb = XGBRegressor()

param_grid = {
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7],
    'min_child_weight': [1, 3, 5],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0]
}

grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid,
                           scoring='neg_mean_squared_error', cv=5, verbose=1)

grid_search.fit(x_train, y_train)

print("Best Parameters:", grid_search.best_params_)
print("Best CV Score:", grid_search.best_score_)

best_xgb = grid_search.best_estimator_

pred_xgb_tuned = best_xgb.predict(x_test)
```

```
mae_xgb_tuned = mean_absolute_error(y_test, pred_xgb_tuned)
mse_xgb_tuned = mean_squared_error(y_test, pred_xgb_tuned)
rmse_xgb_tuned = np.sqrt(mse_xgb_tuned)
rsq_xgb_tuned = r2_score(y_test, pred_xgb_tuned)
```

```
print("\nTuned Model Metrics:")
print("MAE: %.3f" % mae_xgb_tuned)
print("MSE: %.3f" % mse_xgb_tuned)
print("RMSE: %.3f" % rmse_xgb_tuned)
print("R-Squared: %.3f" % rsq_xgb_tuned)
print("Training Accuracy:", best_xgb.score(x_train, y_train))
print("Testing Accuracy:", best_xgb.score(x_test, y_test))
```

Fitting 5 folds for each of 243 candidates, totalling 1215 fits  
Best Parameters: {'colsample\_bytree': 0.8, 'learning\_rate': 0.1, 'max\_depth': 3, 'min\_child\_weight': 1, 'subsample': 0.6}  
Best CV Score: -16626.005239377753

Tuned Model Metrics:  
MAE: 94.131  
MSE: 14517.358  
RMSE: 120.488  
R-Squared: 0.991  
Training Accuracy: 0.9951537856788809  
Testing Accuracy: 0.9910156029061967